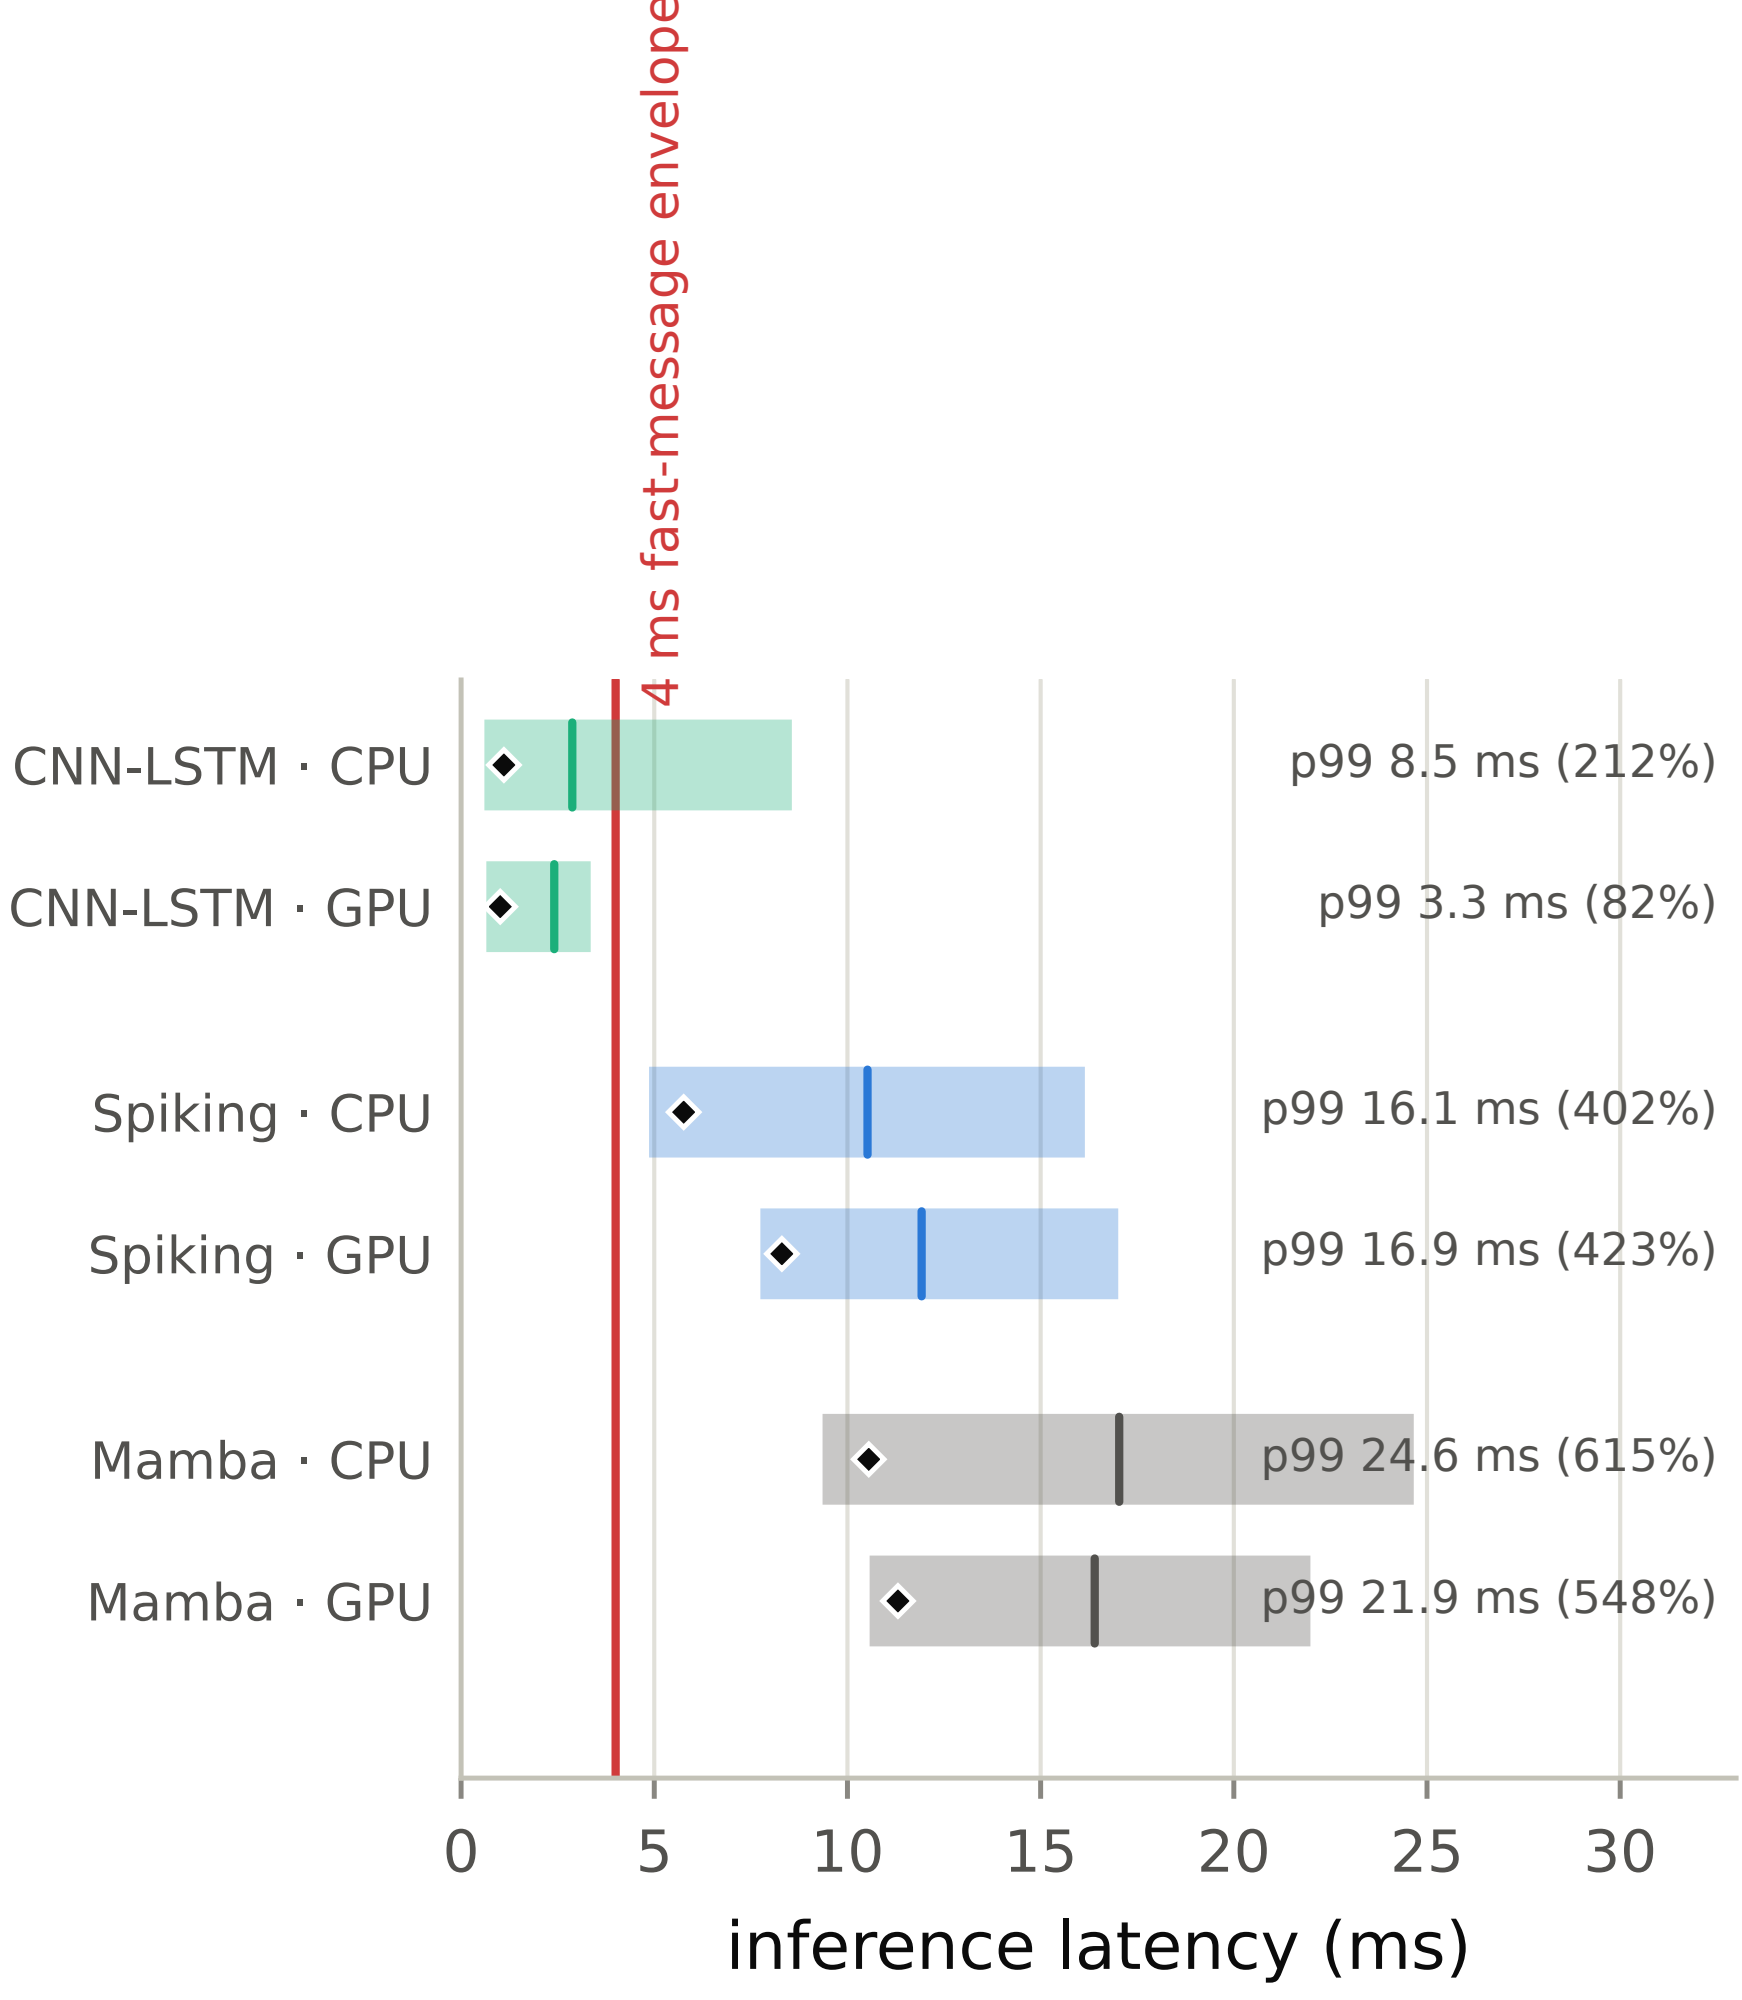
4 ms fast-message envelope
CNN-LSTM · CPU
p99 8.5 ms (212%)
CNN-LSTM · GPU
p99 3.3 ms (82%)
Spiking · CPU
p99 16.1 ms (402%)
Spiking · GPU
p99 16.9 ms (423%)
Mamba · CPU
p99 24.6 ms (615%)
Mamba · GPU
p99 21.9 ms (548%)
0
5
10
15
20
25
30
inference latency (ms)